

# Machine Learning

## (機器學習)

### Lecture 15: Machine Learning Soundings

Hsuan-Tien Lin (林軒田)

`htlin@csie.ntu.edu.tw`

Department of Computer Science  
& Information Engineering

National Taiwan University  
(國立台灣大學資訊工程系)



# Roadmap

- 1 When Can Machines Learn?
- 2 Why Can Machines Learn?
- 3 How Can Machines Learn?
- 4 How Can Machines Learn Better?
- 5 Embedding Numerous Features: Kernel Models
- 6 Combining Predictive Features: Aggregation Models
- 7 Distilling Implicit Features: Extraction Models

## Lecture 15: Machine Learning Soundings

- Deep Learning Initialization
- Deep Learning Optimization
- Deep Learning Regularization

# Deep Learning Initialization

# Weight Initialization

- all 0: too symmetric for tanh, not differentiable for ReLU
- constants: 'cloned' neurons
- too large: saturation/gradient vanishing for tanh, half dying for ReLU

## want

- random: avoid all-0 or constants
- small: 'well-behaved' initialization

next: small random initialization with  
zero-mean (easier to analyze)

# Small Random Initialization for Forwarding tanh

$$\text{score } \mathbf{s}_j^{(\ell)} = \sum_{i=0}^{d^{(\ell-1)}} \mathbf{w}_{ij}^{(\ell)} \cdot \mathbf{x}_i^{(\ell-1)}, \text{ transformed } \mathbf{x}_j^{(\ell)} = \phi_j^{(\ell)} \left( \mathbf{s}_j^{(\ell)} \right)$$

- $w_{ij}^{(\ell)}$  small  $\Rightarrow s_j^{(\ell)}$  small  $\Rightarrow \tanh'(s_j^{(\ell)}) \approx 1$  (approximately linear)

$$\begin{aligned} \text{var}(x_j^{(\ell)}) &\approx \text{var}(s_j^{(\ell)}) = \text{var} \left( \sum_{i=0}^{d^{(\ell-1)}} w_{ij}^{(\ell)} \cdot x_i^{(\ell-1)} \right) \\ &= \text{var}(w_{0j}^{(\ell)}) + \sum_{i=1}^{d^{(\ell-1)}} \text{var}(w_{ij}^{(\ell)} \cdot x_i^{(\ell-1)}), \text{ by independence} \end{aligned}$$

- (ind.):  $\text{var}(wx) = E(x)^2 \text{var}(w) + E(w)^2 \text{var}(x) + \text{var}(w) \text{var}(x)$ 
  - $E(w) = 0$  by construction  $\Rightarrow E(x_i^{(\ell)}) = 0$
  - $E(x_i^{(0)}) = 0$  for mean-shifted features

'ideal'  $\text{var}(w) = 1/d^{(\ell-1)}$  so

$$\text{var}(x_j^{(\ell)}) \approx \text{var}(x_i^{(\ell-1)})$$

# Small Random Initialization for Backward tanh

$$\delta_j^{(\ell-1)} = \sum_k \left( \delta_k^{(\ell)} \right) \left( w_{jk}^{(\ell)} \right) \left( \phi' \left( s_j^{(\ell-1)} \right) \right)$$

- assume approximately linear:  $\phi' \approx 1$
- to keep  $\text{var}(\delta_j^{(\ell-1)})$  similar to  $\text{var}(\delta_k^{(\ell)})$ :

$$\text{var}(w) = \frac{1}{d^{(\ell)}}$$

Xavier initialization (Xavier Glorot, 2010): let

$$\text{var}(w) = \frac{2}{d^{(\ell-1)} + d^{(\ell)}}$$

# Small Random Initialization for ReLU

$$\begin{aligned} \text{var}(s_j^{(\ell)}) &= \text{var} \left( \sum_{i=0}^{d^{(\ell-1)}} w_{ij}^{(\ell)} \cdot x_i^{(\ell-1)} \right) \\ &= \text{var}(w_{0j}^{(\ell)}) + \sum_{i=1}^{d^{(\ell-1)}} \text{var}(w_{ij}^{(\ell)} \cdot x_i^{(\ell-1)}), \quad \text{by independence} \end{aligned}$$

- (ind.):  $\text{var}(wx) = E(x)^2 \text{var}(w) + E(w)^2 \text{var}(x) + \text{var}(w) \text{var}(x)$ 
  - $E(w) = 0$  by construction
  - $\text{var}(wx) = \text{var}(w) E(x^2)$
  - $E(x^2) = 0.5 \text{var}(s^2)$  because of 'ReLU'

He initialization (Kaiming He, 2015):  
 $\text{var}(w) = 2/d^{(\ell-1)}$  so  $\text{var}(s_j^{(\ell)}) \approx \text{var}(s_i^{(\ell-1)})$

**Questions?**



# Deep Learning Optimization

# Difficulty of Deep Learning Optimization

## error surface complicated

- local minima: not as bad as imagined
- saddle points/local maxima: easily escapable (especially with SGD)
- plateau: need larger learning rate  $\eta$
- ravines: need to avoid oscillation

## stability <> computation trade-off

slow computation of gradient (backprop)

⇒ SGD on minibatch

⇒ 'instable' estimate of gradient

getting more stable estimate? **averaging**

# Running Average Estimate of Gradient

gradient descent:  $\mathbf{w}_t \leftarrow \mathbf{w}_{t-1} - \eta \cdot \mathbf{v}_t$

## original minibatch SG

gradient estimate  $\mathbf{v}_t = \Delta_t$  from one minibatch SG

## averaging by multiple SG

if minibatch SG for  $M$  times at  $t$ -th iteration, each getting  $\Delta_t^{(m)}$ , more stable gradient estimate by uniform averaging  $\mathbf{v}_t = \frac{1}{M} \sum_{m=1}^M \Delta_t^{(m)}$  —needing  $M$  times more computation than original minibatch SGD

speedup by reusing each  $\Delta_{t-m+1}$  as  $\Delta_t^{(m)}$

$\mathbf{v}_t = \frac{1}{M} \sum_{m=1}^M \Delta_{t-m+1}$  —‘moving window’ average of SG

issue with ‘moving window’ average:

**uniformly weighted**

# Averaging SG Non-uniformly

## Running Average

- $$\mathbf{v}_t = \beta \mathbf{v}_{t-1} + (1 - \beta) \Delta_t$$
with  $0 \leq \beta < 1$  to control how much history to take
- $\beta = 0$ : original SGD

$$\mathbf{v}_t = \sum_{m=1}^t \beta^{t-m} (1 - \beta) \Delta_t$$

—size- $t$  window, exponentially-decreasing averaging

SGD **with momentum**: optimization direction  
= current SG ( $\Delta_t$ ) + historical inertia ( $\mathbf{v}_{t-1}$ )

# Benefits of SGD with Momentum

$$\mathbf{v}_t = \beta \mathbf{v}_{t-1} + (1 - \beta) \Delta_t$$

$$\mathbf{w}_t = \mathbf{w}_{t-1} - \eta \mathbf{v}_t$$

- some variance in SG canceled out
- oscillation across ravine dampened
- shallow local optima/saddle points escaped

SGD with momentum: 'stabilize' SG with running average

# Per-Component Learning Rate

fixed learning rate :  $\mathbf{w}_t = \mathbf{w}_{t-1} - \eta \mathbf{v}_t$

per-component learning rate :  $\mathbf{w}_t = \mathbf{w}_{t-1} - \eta_t \odot \mathbf{v}_t$

intuition: scales error surface

want: smaller step for larger gradient  
component

# Running Average of Gradient Magnitude

want: smaller step for larger gradient component, say

$$\eta_t = \frac{1}{\nabla E(\mathbf{w}_t) \odot \nabla E(\mathbf{w}_t)} \text{ per component}$$

- full gradient  $\nabla E$  not available, SG only
- using  $\Delta \odot \Delta$  directly: not very stable

idea: running average of  $\Delta_t \odot \Delta_t$

# RMSProp

$$\mathbf{u}_t = \beta \mathbf{u}_{t-1} + (1 - \beta) \Delta_t \odot \Delta_t$$

$$\eta_t = \frac{\eta}{\sqrt{\mathbf{u}_t + \epsilon}} \text{ per component}$$

$$\mathbf{w}_t = \mathbf{w}_{t-1} - \eta_t \odot \Delta_t$$

RMSProp: SGD + per-component learning rate  
using running average of magnitude



# Adam: Adaptive Moment Estimation

Adam  $\approx$  momentum + RMSProp + global decay

$$\mathbf{v}_t = \beta_1 \mathbf{v}_{t-1} + (1 - \beta_1) \Delta_t$$

$$\mathbf{u}_t = \beta_2 \mathbf{u}_{t-1} + (1 - \beta_2) \Delta_t \odot \Delta_t$$

$$\eta_t = \frac{\eta}{\sqrt{\mathbf{u}_t + \epsilon}} \cdot \frac{1}{\sqrt{t/N}}$$

$$\mathbf{w}_t = \mathbf{w}_{t-1} - \eta_t \odot \mathbf{v}_t$$

- momentum in  $\mathbf{v}_t$
- RMSProp in  $\mathbf{u}_t$
- global decay by  $\sqrt{t/N}$
- (some minor correction of estimation)


Adam usually more aggressive than original SGD (but can also overfit faster)

**Questions?**


# Deep Learning Regularization

# A Basic Trick: Early Stopping


- **GD/SGD (backprop)** visits more weight combinations as  $t$  increases



- smaller  $t$  effectively decrease  $d_{VC}$
- better 'stop in middle': **early stopping**



( $d_{VC}^*$  in middle, remember? :-))



when to stop? **validation!**

# Co-Adaptation Issue of Deep Learning


## co-adaptation

- consistent **mistakes** from some neurons: like **noise**
  - corrected by **fitting** other neurons: like **overfit**
- unhealthy dependence between neurons

co-adaptation harms generalization of deep learning

# Breaking the Dependence

idea: **shut down** some neurons randomly



**dropout:**

- drop  $p$ , keep  $1 - p$
- implicit aggregation of many **thinner networks**
- slow down convergence, **but faster per-iteration**

**dropout:** simple yet effective technique for deep learning regularization

## Dropout During Testing

## \* dropout during testing

- full-net prediction w/o changing

$$\underset{\text{in training}}{E(S_i^{(l)})} = (1-p) \underset{\text{in testing}}{S_i^{(l)}}$$

- test-time "pseudo-" dropout

$$x_i^{(l)} = \theta(\underbrace{(1-p)}_{\text{need to record } p, \text{ less flexibility}} S_i^{(l)})$$

need to record  $p$ , less flexibility

- inverted dropout

for changing  $p$  per neuron  
or dynamically

training : dropout &  $x_i^{(l)} = \theta(S_i^{(l)} / (1-p))$

testing :  $x_i^{(l)} = \theta(\underbrace{(1-p')}_{\text{unchanged}} S_i^{(l)} / \underbrace{(1-p')}_{\text{unchanged}})$

$$= \theta(S_i^{(l)}) \quad \text{unchanged (usually preferred)}$$

**Questions?**